

After the Launch Party

Flagship Feature Reference

What each flagship feature is, what it added to the product that wasn't available before, and how the product improved because of it.

Purpose of this doc: A companion to the Company Reference. The company doc answers “what is this company?” This doc answers “what is the specific feature we’re scoring, and what changed about the product when it shipped?” You’ll use it during scoring as a sanity check: if the “what changed” story here doesn’t hold up under evidence later, the survival scoring will catch it.

A note on framing: For some features, the “what wasn’t available before” is a clean new capability (Claude Artifacts introduced generative UI rendering). For others, it’s more about degree, positioning, or strategic shift (ChatGPT Memory persisted a capability that already technically existed in context windows). I’ll be honest about that distinction per entry rather than forcing every feature into a dramatic before and after frame.

Version: v1.0, April 2026. Iterate as scoring surfaces corrections.

Group A, AI Native (15 Features)

1. ChatGPT Memory

Launched: February 2024 (research preview), September 2024 (general availability across Free, Plus, Team, Enterprise), April 2025 (major upgrade to reference all past conversations, not just explicitly saved items).

What the feature is: A persistent memory layer in ChatGPT. The AI remembers things you've told it, or picks up on things itself, and carries that context across future conversations. You can view, edit, or delete memories from settings.

What wasn't available before: Pre Memory, every ChatGPT conversation started from scratch. If you told it your dietary preferences, writing style, job context, or preferences, you had to retell it every new chat. The model had no persistent model of you as a user.

How the product improved:

- **Personalization without re prompting.** Users stopped repeating themselves across sessions, which is the largest quality of life improvement most regular ChatGPT users had experienced since the original launch.
- **Cross session consistency.** Writing style, preferences, and context stayed stable across chats, which made ChatGPT feel less like a tool and more like a collaborator.
- **Raised the bar for the whole category.** Within 18 months, Claude, Gemini, and virtually every other assistant shipped their own memory features. Memory became table stakes.
- **Created new product surface.** Memory enabled downstream features like Study Mode and persistent Projects, which wouldn't have worked without a persistent user model.

Strategic significance for the project: Memory is arguably the single most consequential AI feature OpenAI shipped in 2024. Scoring whether it actually changed user behavior (vs. being a toggle people forget exists) is one of the highest stakes questions in the dataset.

2. Claude Artifacts

Launched: June 2024.

What the feature is: A side panel in the Claude chat interface where Claude can generate and render standalone content (code, HTML, React components, SVG diagrams, documents, etc.). The user can interact with the rendered output, iterate on it, and export it, while the chat continues in the main thread.

What wasn't available before: Pre Artifacts, Claude output code and content as text inside the chat. The user had to copy it, paste it into an editor or tool, and see it render there. There was no way to see a live rendered output alongside the conversation, and no clean separation between “discussion” and “artifact.”

How the product improved:

- **Introduced generative UI as a mainstream concept.** Artifacts let users see AI generated code actually run in a sandbox. This turned Claude from a text generator into something closer to a development environment.
- **Collapsed workflow for creators.** Designers, developers, and analysts who used Claude for prototyping now had a single surface for chat plus preview, instead of juggling ChatGPT plus a code sandbox.
- **Created a pattern the whole industry copied.** Within six months, OpenAI shipped Canvas, Google shipped Gemini Canvas, and several smaller labs shipped their own versions. Artifacts became the reference for how chat products should handle structured output.
- **Opened new use cases for Claude.** Interactive diagrams, custom calculators, shareable prototypes, and dashboards became realistic outputs. These were impractical in pure text.

Strategic significance for the project: Artifacts is one of the cleanest examples in the dataset of an AI feature that created language. The interesting scoring question is whether Anthropic held the narrative advantage or lost it to the copycats.

3. Perplexity Deep Research

Launched: February 2025.

What the feature is: A mode where Perplexity runs an extended, multi step research process on a user's question, consulting dozens of sources, reasoning across them, and producing a structured multi page report with citations. Takes 2 to 5 minutes per query rather than seconds.

What wasn't available before: Pre Deep Research, Perplexity's responses were fast cited answers (seconds long) and Pages (user assembled research documents). There was no mode for *Perplexity itself* to conduct an extended research process on the user's behalf. Users who wanted depth had to chain multiple searches manually or use OpenAI's own Deep Research tier.

How the product improved:

- **Turned Perplexity from a search engine into a research assistant.** The product category shifted. Instead of competing with Google, Deep Research put Perplexity in competition with junior analysts.
- **Justified the premium tier.** Deep Research gave Pro and Max subscribers a feature worth the price that wasn't available in free or competitors' basic tiers.
- **Elevated output format.** Reports included structured sections, citation density, and breadth that felt publication quality rather than chat quality.
- **Responded directly to OpenAI's own Deep Research.** The feature ships happened within weeks of each other, making this one of the clearest head to head product moments of 2025.

Strategic significance for the project: Deep Research is the feature that defines whether Perplexity can hold a premium position as frontier labs encroach on search. Scoring its stickiness tells you a lot about Perplexity's long term defensibility.

4. Glean Assistant

Launched: General availability mid 2024 (extensive beta period before).

What the feature is: Glean’s conversational AI layer. Sits on top of Glean’s enterprise search and connects to all the work apps a company uses, lets employees ask natural language questions across that content, and returns grounded answers with citations. Increasingly also acts on workflows.

What wasn’t available before: Pre Assistant, Glean was universal search across company apps. You got a ranked list of results. The AI Assistant turned that into a Q&A layer: instead of “here are 10 documents that mention Q3 budget,” you got a summarized answer.

How the product improved:

- **Closed the loop from search to answer.** Users stopped having to read through the top results to synthesize the answer themselves.
- **Made Glean viable as a primary interface for knowledge work.** Employees could live inside Glean rather than bouncing between apps.
- **Added agent like capabilities.** Later Assistant updates let it take actions in connected apps (draft emails, create Jira tickets, update Salesforce records), which moved Glean toward being an AI workplace surface rather than a search tool.
- **Strengthened competitive positioning vs. Microsoft Copilot and ChatGPT enterprise.** The “AI that knows your company” pitch now had a clean product to point to.

Strategic significance for the project: Glean Assistant is the frontier of “AI for work” in enterprise. Its adoption signal is one of the cleanest proxies for whether enterprises are actually deploying AI beyond pilots.

5. Cursor Composer (Agent Mode)

Launched: November 2024 (Composer), evolved into Agent mode through early 2025.

What the feature is: An agent inside Cursor that can edit multiple files in a codebase based on a natural language instruction. You describe what you want (“refactor the auth module to use JWT,” “add dark mode across all components”), and Cursor’s agent plans the edits, applies them across files, runs commands if needed, and presents a diff for you to accept or reject.

What wasn’t available before: Pre Composer, Cursor’s AI was primarily autocomplete (Cursor Tab) and chat based edits to a single file. For changes spanning multiple files, you either guided it file by file or used inline chat in each location. There was no holistic “agent that reads your whole codebase and changes many files at once.”

How the product improved:

- **Shifted coding AI from assistance to delegation.** Developers moved from “AI helps me write this line” to “AI, handle this whole feature.”
- **Made Cursor materially more valuable than GitHub Copilot.** Copilot was (and largely remains) line level. Composer is a different category entirely.
- **Unlocked new developer use cases.** Large scale refactors, migrations, test suite updates, and first drafts of full features became routine agent tasks rather than hours of manual work.
- **Forced the whole AI coding category to respond.** GitHub shipped Copilot Workspace, Replit upgraded Replit Agent, Anthropic shipped Claude Code, and Windsurf appeared. Composer set the reference.

Strategic significance for the project: Composer is one of the cleanest “AI feature that clearly won” stories in the dataset. Its survival signal is strong (retention, public enthusiasm, enterprise deployment) and serves as the benchmark against which other features compare.

6. Harvey Assistant

Launched: 2023 general availability, major 2024 expansion into a broader suite.

What the feature is: Harvey’s core LLM powered legal assistant. Lawyers use it to research cases, draft documents, analyze contracts, run diligence workflows, and answer legal questions grounded in firm specific and jurisdiction specific sources.

What wasn’t available before: Pre Harvey, legal research and drafting at large firms used tools like LexisNexis, Westlaw, or iManage, which were essentially sophisticated search engines. Lawyers did the synthesis and drafting themselves. The AI layer that could actually *draft* legal content in a firm’s voice and style didn’t exist.

How the product improved:

- **Compressed the junior associate’s repetitive work.** First draft memos, contract markups, and diligence summaries dropped from hours to minutes. Law firms adopted Harvey faster than any AI in legal history.
- **Made “AI in legal” credible.** Before Harvey, most legal AI failed the accuracy bar or the confidentiality bar. Harvey cleared both at most of the top 100 firms.
- **Created a reference for vertical AI.** Harvey became the canonical case study for why domain specific AI could outperform general purpose (ChatGPT) in high stakes professional work.
- **Expanded Harvey’s strategic scope.** From a research assistant, Harvey grew into workflow specific products (Harvey Workflows) that let firms automate repeatable matter types.

Strategic significance for the project: Harvey is the dataset’s vertical AI anchor. Scoring its stickiness tests whether narrow, domain specific AI holds its advantage as general models get better.

7. Jasper AI Agents Platform

Launched: Rolled out through 2024, fully platform branded in 2025 with 100+ specialized marketing agents.

What the feature is: Jasper’s pivot from “AI writing tool” to “AI agent workspace for marketing teams.” The platform includes agents for specific marketing tasks (blog drafting, social post generation, ad copy, SEO optimization, campaign planning, localization), structured content pipelines that chain agents into end to end workflows, and Brand Voice to keep output consistent.

What wasn’t available before: Pre Agents, Jasper was a document editor with AI writing assistance and templates. You wrote a prompt, got a draft, edited it. There was no concept of handing off a multi step marketing task and having the system plan and execute it across channels.

How the product improved:

- **Positioned Jasper as an agent platform, not a writing tool.** This was a strategic reframe designed to counter the fact that ChatGPT had commoditized basic AI writing.
- **Enabled end to end campaign workflows.** Briefs in, assets across blog, email, social, ads out. One system instead of a dozen apps.
- **Deepened brand enforcement.** Brand Voice got more sophisticated and was embedded across all agents, addressing the single biggest enterprise concern about AI writing: consistency.
- **Opened enterprise pricing.** The agent platform justified higher tier pricing, which was necessary as the “AI writing tool” pricing ceiling dropped.

Strategic significance for the project: Jasper is the most important *adaptation story* in the dataset. Whether the pivot worked determines whether first wave AI companies can survive commoditization.

8. Writer AI Studio

Launched: 2024.

What the feature is: Writer’s platform for enterprises to build custom AI workflows, grounded on their own content via Knowledge Graph. Non technical operators can assemble AI apps for specific use cases (marketing briefs, legal review, HR responses) without engineering support.

What wasn’t available before: Pre AI Studio, Writer was primarily a writing assistant with some enterprise deployment features. There was no path for non engineers at a Fortune 500 to build custom domain specific AI apps without writing code or calling LLM APIs directly.

How the product improved:

- **Made Writer a platform play, not a single product.** Enterprise customers got a system they could invest in long term, not just a writing tool.
- **Strengthened Writer’s differentiation against general AI.** “Build your own AI apps grounded on your own content” is a pitch that works against ChatGPT Enterprise because it addresses IT’s biggest concern (data boundary).
- **Enabled direct ROI stories.** Because AI Studio apps are built for specific measurable tasks, Writer could point customers at concrete outcomes (hours saved, tickets deflected).
- **Opened larger contract sizes.** Platform pricing with custom apps is meaningfully higher ACV than per seat writing tools.

Strategic significance for the project: Writer is the dataset’s weakest entry due to thin public evidence (enterprise, sales led go to market). AI Studio is the most defensible flagship pick because it’s the strategic bet that matters, even if scoring it will rely more on press and earnings than user reviews.

9. Gamma AI Generator (Text to Deck)

Launched: Core feature since 2022 launch; major LLM upgrades through 2024 including Gamma 2.0.

What the feature is: Gamma’s core AI interaction. You type a description or paste content (“a pitch deck for a Series A fintech focused on cross border payments”), and Gamma generates a fully designed, editable presentation with slides, images, layouts, and content. You can iterate with further prompts or edit manually.

What wasn’t available before: Pre Gamma, creating a presentation meant either starting from a blank PowerPoint/Keynote/Google Slides file, or using a template library (Canva, Slidebean). There was no “describe it and get a designed deck” flow. Even early AI design tools generated text, not structured, formatted presentations.

How the product improved:

- **Made “describe, get deck” a real product pattern.** Gamma proved that AI could generate not just content but also design decisions (layout, visual hierarchy, imagery).
- **Compressed deck creation from hours to minutes for most use cases.** First drafts for internal decks, investor updates, and class presentations stopped being the bottleneck.
- **Expanded the product category.** Gamma extended beyond decks into documents and websites, but text to deck remains the flagship interaction and the reason users sign up.
- **Shifted user expectations of AI and design tools.** Canva, Microsoft, and Google all followed with their own AI presentation tools, but Gamma holds the “native AI deck” category.

Strategic significance for the project: Gamma’s text to deck is one of the cleanest “AI native category creators” in the dataset. The scoring question is whether Gamma’s early lead translated into lasting market position, or whether Canva’s scale and Microsoft’s distribution caught up.

10. Descript Underlord

Launched: 2024.

What the feature is: Descript's named AI agent that sits alongside the editing interface. It can take high level instructions ("cut out filler words, smooth out the pacing, find the best 60 seconds for a social clip, write a YouTube description") and execute them across a video or audio project. Underlord coordinates Descript's existing AI tools (Studio Sound, Overdub, transcript editing) plus new agent level actions.

What wasn't available before: Pre Underlord, Descript had many individual AI features (Overdub, Studio Sound, filler word removal, Eye Contact), but the user had to know which tool to use and trigger each manually. There was no "ask an AI to handle this for me" agent layer.

How the product improved:

- **Lowered the skill floor.** Non editors could now get professional results without learning which tool does what.
- **Enabled multi step workflows.** "Make this a LinkedIn video" became one instruction instead of a sequence of discrete operations.
- **Positioned Descript against Capcut and Runway.** Both competitors had been pushing AI native editing. Underlord gave Descript a named, marketable AI agent to compete on.
- **Unlocked new use cases.** Batch processing of podcast episodes, automatic social clip generation, and templated content pipelines became viable.

Strategic significance for the project: Underlord is the feature that defines whether Descript stays relevant as AI native video tools multiply. Its survival signal is a meaningful proxy for incumbent creative tools' ability to adapt.

11. Runway Gen-3 Alpha

Launched: June 2024.

What the feature is: Runway’s text to video and image to video model. Generates short (up to 10 second) video clips from text prompts or starting images, with photorealistic quality, camera control, and motion dynamics far beyond previous generations.

What wasn’t available before: Pre Gen-3, AI video was either short, low quality, or both (Gen-2 produced 4 second clips with visible artifacts). OpenAI had announced Sora in February 2024 with stunning demos but no availability to users. Gen-3 was the first production grade video generation model shipped to paying users.

How the product improved:

- **Made AI video production viable for real workflows.** Ad agencies, music video creators, and filmmakers started using Runway in commercial production, not just experimentation.
- **Set a quality bar that Runway defended through 2025.** Even as Luma, Kling, Pika, and eventually Sora launched, Gen-3 and its successors kept Runway in the top tier.
- **Expanded the creative toolkit.** Motion Brush, Act One (performance capture), and camera controls were layered on top of Gen 3, making it more than a one shot generator.
- **Created the head to head Sora comparison.** The Runway vs. Sora story dominated AI video coverage for months; Runway won the “ships first, iterates faster” argument.

Strategic significance for the project: Gen-3 is one of the dataset’s clearest “AI native category leader” stories. Scoring it tests whether shipping first translated to durable advantage, or whether OpenAI’s distribution eventually won.

12. ElevenLabs Conversational AI

Launched: Late 2024.

What the feature is: A platform for building real time conversational voice agents. Developers can configure an agent with a voice, personality, and knowledge base, and deploy it into phone calls, apps, games, or support workflows. Handles turn taking, interruptions, and multi language conversations.

What wasn't available before: Pre Conversational AI, ElevenLabs was a text to speech and voice cloning engine. Developers used it to generate pre rendered audio (for audiobooks, videos, game characters). There was no native way to build a real time conversational agent using ElevenLabs voices. Developers had to stitch together ElevenLabs plus an LLM plus a speech to text service plus orchestration logic themselves.

How the product improved:

- **Repositioned ElevenLabs from component vendor to platform.** Instead of selling voices, they now sell the whole voice agent stack.
- **Entered the voice agent wave directly.** As OpenAI launched Realtime API and startups shipped voice first products, ElevenLabs had an integrated offering to compete.
- **Enabled new customer segments.** Customer support, sales, and consumer app companies that wanted voice agents without assembling the stack now had a one stop product.
- **Defended the core voice tech moat.** By building the platform on their own voices, ElevenLabs kept customers using their TTS at a higher price point than raw API calls.

Strategic significance for the project: Conversational AI is ElevenLabs' bet to avoid commoditization as frontier labs add competitive voice models. Its adoption is a clean proxy for whether voice AI is a real enterprise wave or a feature category that gets absorbed.

13. Character.ai (Core Character Chat Experience)

Launched: 2022 (core product), with continuous iteration; major inflection in late 2024 with Google founder departure and product pivots.

What the feature is: The core Character.ai experience: users create or select an AI character with a defined personality, backstory, and voice, then have ongoing text (and sometimes voice) conversations with it. Characters range from fictional (anime characters, historical figures, celebrities) to original creations. Conversations persist, and users return to the same characters over weeks or months.

What wasn't available before: Pre Character.ai, AI chat was largely single purpose (ChatGPT for tasks, Replika for companionship). Character.ai introduced at massive scale the idea that users would form ongoing relationships with *multiple* AI characters across entertainment, learning, and emotional support use cases. Session times reached multi hour averages, higher than any consumer AI product.

How the product improved (then regressed):

- **Pioneered consumer AI companionship at scale.** Character.ai was the dominant consumer AI app by session time in 2023 to early 2024.
- **Proved a UGC model for AI.** Users creating characters for other users created network effects that no other AI product has replicated.
- **Then faced serious crises.** Lawsuits over minor user harm, founder departure to Google, and user backlash over model downgrades (users said the “new” characters felt duller and more restricted) created a visible decline narrative through 2024 to 2025.

Strategic significance for the project: Character.ai is the dataset's consumer AI companion anchor and its clearest cautionary tale. Scoring it honestly (measuring, not moralizing) tests whether consumer AI companionship has a durable business, or whether its early traction was peak and the decline is terminal.

14. Replit Agent

Launched: September 2024.

What the feature is: An autonomous coding agent inside Replit. You describe an application in natural language (“build me a recipe app with user accounts and photo uploads”), and Replit Agent plans the build, writes the code across files, sets up the database, configures deployment, and produces a working app. Non engineers can build real applications; engineers use it for rapid prototyping.

What wasn’t available before: Pre Agent, Replit was a browser based coding environment with AI assistance (Ghostwriter). Developers wrote code with AI help. There was no flow where a non coder could describe an app in English and end up with a deployed, working application. Cursor’s Composer came later in 2024 and targeted working developers; Replit Agent was the first mainstream bet that AI could build apps for non engineers.

How the product improved:

- **Opened Replit to a new audience.** Non engineers (founders, designers, students, hobbyists) started building real apps on Replit.
- **Shortened the “idea to deployed app” loop dramatically.** Prototypes that took days now took hours.
- **Redefined Replit’s market position.** From a coding environment for developers to a “build apps with AI” platform, competing directly with Lovable, Bolt, v0, and Cursor.
- **Forced a business model shift.** Replit’s pricing, hosting, and packaging all had to adapt to agent style usage with very different compute patterns.

Strategic significance for the project: Replit Agent is the flagship “AI builds software” bet of the dataset. Its survival signal tells you whether non engineer app building is a durable category or a hype cycle. Highly consequential for both Replit and the broader “AI coding” space.

15. Granola Spaces

Launched: May 2025 (Granola 2.0 release).

What the feature is: Team workspaces in Granola. Users can create shared folders (Sales Calls, Customer Feedback, Hiring Loops, Weekly Syncs), with Granola storing meeting notes from team members into the appropriate Space. Users can Chat with Folders to query across an entire Space (“why are we losing deals this quarter?”) and get answers grounded in all the relevant meeting notes.

What wasn’t available before: Pre Spaces, Granola was an individual meeting notepad. Each user’s notes were personal. Teams that wanted shared meeting intelligence had to manually share notes in Slack or Notion, and there was no cross meeting AI query across a team’s conversations. Granola’s AI was scoped to your own meetings only.

How the product improved:

- **Turned Granola from a prosumer app into an enterprise context layer.** The pitch shifted from “transcribe your meetings” to “turn your team’s meetings into organizational memory.”
- **Unlocked cross meeting insights.** Product, sales, and recruiting leaders could ask questions that spanned dozens of meetings and get synthesized answers.
- **Justified enterprise pricing and contracts.** Spaces plus SSO, admin controls, and Slack integration created a real enterprise product that could command contracts vs. per seat individual subscriptions.
- **Supported the \$125M Series C.** Granola’s March 2026 raise at a \$1.5B valuation was pitched on Spaces’ traction and the enterprise pivot. Without Spaces, the story is different.

Strategic significance for the project: Spaces is Granola’s strategic bet on moving from commoditized individual meeting notes to defensible enterprise context infrastructure. Its adoption is the single most important signal for whether Granola earns its valuation.

Group B, Traditional SaaS (15 Features)

16. Notion Q&A

Launched: November 2023.

What the feature is: An AI powered workspace search and Q&A feature. Users ask natural language questions (“what did we decide about the Q3 pricing change?” or “what’s our current hiring plan?”), and Notion searches across all pages in the workspace, synthesizes an answer, and cites the source pages.

What wasn’t available before: Pre Q&A, Notion had keyword search. You had to remember which page something was on, type the right words, and open documents to find the actual information. There was no way to ask a question and get an answer grounded in your workspace content.

How the product improved:

- **Turned Notion from “where knowledge is stored” into “where knowledge is used.”** The retrieval problem (finding what you already wrote) was the single biggest pain point for Notion power users. Q&A addressed it directly.
- **Positioned Notion against Glean and Dropbox Dash.** Notion users no longer needed a separate tool for workspace AI search if Q&A worked.
- **Justified the Notion AI premium tier.** Q&A was a premium add on, which meant Notion had a named feature customers would pay extra for.
- **Set a pattern for workspace AI.** Many productivity tools followed with their own “ask your workspace” features.

Strategic significance for the project: Q&A is the most important AI feature Notion shipped, and whether it actually gets used (vs. users going to Glean or ChatGPT for workspace search) is a high stakes strategic question Notion has not answered publicly.

17. Slack AI

Launched: February 2024.

What the feature is: A bundle of AI features inside Slack: channel recap (summaries of what you missed), thread summaries, AI powered search with natural language answers, and daily recaps across your channels. Sold as a premium add on for paid Slack workspaces.

What wasn't available before: Pre Slack AI, catching up on Slack was manual scrolling. Users either read every message, used channel notifications heuristics, or gave up and asked colleagues what they missed. Search was keyword based, returning messages rather than answers.

How the product improved:

- **Addressed Slack's biggest user complaint.** Information overload in channels was the most cited frustration in user research for years. Recap directly addressed it.
- **Monetized AI in Slack.** The paid add on model made Slack AI a clean revenue signal, unusual in a market where most AI features are bundled free.
- **Fed into Salesforce's Agentforce strategy.** Slack AI became a delivery surface for Agentforce agents that act across Salesforce plus Slack.
- **Strengthened Slack's case against Teams.** Microsoft Teams integrated Copilot; Slack needed its own named AI to compete on sales calls.

Strategic significance for the project: Slack AI is one of the few AI features with a clean paid tier pricing signal, which makes adoption unusually measurable. Scoring whether customers actually pay for it is a key test of “do enterprise users pay for AI or expect it bundled?”

18. Asana AI Studio

Launched: Late 2024.

What the feature is: A no code builder for AI workflows inside Asana.

Operations and project managers can create custom “smart workflows” where AI agents handle specific tasks (triage incoming requests, generate project status, auto assign work based on context, draft responses). Sits alongside older Asana Intelligence features like smart status and smart goals.

What wasn’t available before: Pre AI Studio, Asana had individual AI features (smart goals, smart summaries, smart fields) but no way for customers to build their own AI driven automation. Teams who wanted custom AI automation had to use Zapier plus OpenAI plus glue code, outside of Asana.

How the product improved:

- **Turned Asana into an automation platform.** Customers could automate real workflows instead of just running reports.
- **Addressed a competitive gap.** Monday, ClickUp, and Notion were all shipping AI automation; Asana needed a platform answer.
- **Opened higher tier pricing.** AI Studio is gated to enterprise tiers, which justified existing pricing and supported expansion.
- **Strengthened Asana’s enterprise position.** Larger customers with complex workflows had a reason to consolidate on Asana rather than using it as one of several tools.

Strategic significance for the project: AI Studio is the feature Asana is betting on to compete in the AI work management category. Public evidence is thin, which makes it one of the project’s more uncertain scoring targets.

19. Linear Agents

Launched: 2025.

What the feature is: AI agents that act on Linear issues. They can triage new issues, draft responses, identify duplicates, suggest assignees, escalate blockers, and perform multi step work within Linear’s issue tracking and project planning. Designed in Linear’s characteristic minimalist style, with careful integration and deliberate restraint vs. shipping every possible AI feature.

What wasn’t available before: Pre Agents, Linear had smaller AI features (auto generated issue descriptions, AI triage suggestions, smart search). There was no agent that could *act* on issues at the product level. Teams doing agent style automation used Zapier plus external AI, outside Linear.

How the product improved:

- **Added the agent layer without breaking Linear’s aesthetic.** Linear is known for obsessive design discipline; shipping agents without cluttering the UI was itself a product accomplishment.
- **Strengthened Linear against Jira and Asana.** Both had shipped AI earlier but less coherently. Linear’s “AI should reduce friction, not add buttons” positioning reads as more mature.
- **Retained the brand premium.** Linear’s pricing and positioning depend on feeling more refined than alternatives. Agents had to ship in a way that reinforced rather than diluted that.
- **Opened new use cases.** Triage, duplicate detection, and first pass assignment free engineering managers from repetitive work.

Strategic significance for the project: Linear Agents is an interesting counterpoint in the dataset, a deliberate “ship less, ship better” approach to AI. Scoring whether disciplined AI outperforms aggressive AI shipping is one of the project’s sharper comparisons.

20. Figma AI (Make Designs)

Launched: June 2024, with a public pause of Make Designs shortly after due to quality issues.

What the feature is: Figma’s suite of AI features, announced at Config 2024. Centerpiece: Make Designs, which generated complete UI designs from text prompts. Supporting features: visual search, layer renaming, auto layout suggestions, and AI assisted prototyping. Make Designs was paused when users demonstrated outputs closely resembling existing Apple apps.

What wasn’t available before: Pre Figma AI, designers built UIs manually using components, libraries, and templates. There was no “describe a UI, get a designed screen” flow inside Figma. Competitors like Uizard had shipped AI design tools, but Figma’s scale made this the highest profile AI UI generation launch.

How the product improved (and regressed):

- **Attempted to bring AI UI generation into the dominant design tool.** Before the pause, Make Designs generated entire screens from prompts inside Figma, lowering the skill floor for UI creation.
- **Was met with user backlash.** Users shared outputs that looked near identical to Apple’s Weather app and other shipped products, raising training data questions.
- **Was publicly paused.** Figma acknowledged the issue and paused Make Designs, a rare public retreat from a launched AI feature.
- **Other Figma AI features (search, renaming, prototyping assists) continued.** The narrative damage, though, was tied to Make Designs specifically.

Strategic significance for the project: Figma AI is one of the dataset’s richest stories. A major public AI feature launch with a documented retreat, and an unanswered question about how Figma recovered (or didn’t) in the following year. High insight room and high narrative weight.

21. Canva Magic Studio

Launched: October 2023.

What the feature is: Canva's umbrella AI product, bundling Magic Write (AI copy), Magic Design (generate designs from prompts), Magic Media (AI image, video, graphic generation), Magic Edit (AI photo editing), Magic Switch (resize and reformat across formats), and more. Shipped to all 200M plus Canva users, with some features free and others behind Canva Pro.

What wasn't available before: Pre Magic Studio, Canva was a manual design tool with templates. Users created content by picking a template and editing it. There was no AI generation surface in Canva. Users who wanted AI image generation or writing went to separate tools (Midjourney, ChatGPT) and pasted outputs back into Canva.

How the product improved:

- **Brought generative AI to 200M plus non technical users at once.** The largest consumer AI rollout of the wave by raw audience size.
- **Made Canva Pro more compelling.** Pro unlocked higher Magic Studio usage limits, which pushed free to paid conversions.
- **Addressed the category threat from AI native design tools.** Tome, Gamma, and Beautiful.ai were eating into Canva's deck use case; Magic Design was Canva's answer.
- **Leveraged Canva's scale into AI network effects.** Every Canva user became a trainer/tester, which gave Canva iteration speed advantages.

Strategic significance for the project: Magic Studio is the dataset's largest consumer AI rollout by user count. Scoring whether non technical users actually adopted and retained AI features (vs. clicking once and reverting to templates) is a key cross modality comparison.

22. HubSpot Breeze

Launched: September 2024 (branded umbrella for Breeze Copilot, Breeze Agents, Breeze Intelligence).

What the feature is: HubSpot’s rebranded AI layer. Breeze Copilot is an in-product assistant across HubSpot’s apps. Breeze Agents are specialized AI agents (content, prospecting, social, customer service). Breeze Intelligence provides enriched buyer and company data.

What wasn’t available before: Pre Breeze, HubSpot had many individual AI features (content assistant, ChatSpot, AI insights) scattered across products with inconsistent branding and unclear relationships. There was no coherent “this is HubSpot’s AI” product story. Breeze consolidated the messaging and added the agent layer as a new bet.

How the product improved:

- **Gave HubSpot a clean AI story to sell.** Customers and press could now point to one brand (“Breeze”) instead of parsing a dozen feature names.
- **Shipped genuinely new capabilities.** Breeze Agents (especially the prospecting and customer service agents) were new, not just rebrandings.
- **Positioned HubSpot in the agent wave.** As Salesforce shipped Agentforce and startups shipped specialized agents, HubSpot had a competitive answer at the SMB/mid market layer.
- **Strengthened retention on higher tiers.** Breeze is gated across HubSpot tiers, with the most powerful capabilities requiring higher plans.

Strategic significance for the project: Breeze is HubSpot’s biggest AI brand consolidation and strategic bet. Scoring it tests whether the “AI for SMBs” thesis produces real adoption, or whether SMBs default to general tools (ChatGPT) over embedded CRM AI.

23. Salesforce Agentforce

Launched: September 2024 (Dreamforce announcement), with continuous expansion through 2025 to 2026.

What the feature is: Salesforce’s AI agent platform. Customers can deploy pre built agents (for sales development, customer service, commerce, marketing) or build custom agents on top of Salesforce data, grounded in company specific knowledge and workflows. Agents take actions across Salesforce and connected systems. Priced per conversation.

What wasn’t available before: Pre Agentforce, Salesforce had Einstein (predictions, recommendations) and Einstein GPT/Copilot (generative AI inside Salesforce apps). There was no agent platform, no per conversation pricing, and no pre built vertical agents at this scale. Agentforce represented a major category bet distinct from previous Einstein products.

How the product improved:

- **Shifted Salesforce’s AI narrative from embedded features to autonomous agents.** This reframed Salesforce in investor and customer conversations as an “AI platform” rather than a CRM with AI features.
- **Created a new pricing model (per conversation/per action).** Which both aligned costs to value and opened usage based expansion revenue.
- **Positioned Salesforce against Microsoft, ServiceNow, and standalone agents.** The “your AI agents should live where your data lives” pitch became the explicit competitive wedge.
- **Became the center of Salesforce’s go to market.** Benioff’s earnings calls, Dreamforce keynotes, and Salesforce’s 2025 revenue story all focused on Agentforce adoption.

Strategic significance for the project: Agentforce is the single largest enterprise AI bet on this list by budget and scope. Scoring actual adoption vs. marketing claims is one of the highest stakes and highest insight questions in the entire dataset.

24. Zoom AI Companion

Launched: September 2023 (general availability).

What the feature is: Zoom’s AI assistant across their product suite. Features include meeting summaries, smart recordings with chapters and action items, chat compose, meeting Q&A (ask questions about a past or live meeting), email and calendar drafting. Notably, included free with most paid Zoom plans rather than sold as a premium add on.

What wasn’t available before: Pre AI Companion, Zoom meetings had transcription (Zoom IQ, Otter integration) but no integrated AI summarization, chaptering, or cross meeting Q&A native to the paid product. Users relied on third party AI meeting assistants (Otter, Fireflies, Read, Granola) layered on top.

How the product improved:

- **Integrated meeting AI into the default Zoom experience.** No install, no separate app, no new login.
- **Free bundling was a strategic wedge.** Otter, Fireflies, and others charged for features Zoom gave away, forcing them to differentiate or lose pricing power.
- **Extended across Zoom’s suite (phone, chat, docs).** AI Companion became the thread tying Zoom Workplace together, which supported Zoom’s “beyond meetings” strategy.
- **Strengthened Zoom’s paid tier value.** As competitors (Google Meet, Teams) bundled their own AI, Zoom had to match to keep its base.

Strategic significance for the project: Zoom AI Companion is the dataset’s clearest “distribution driven AI adoption” story. Scoring it tests whether free bundling actually produces sticky usage, or whether users went to dedicated tools anyway.

25. Intercom Fin AI Agent

Launched: March 2023 (Fin v1), 2024 (Fin v2 with more capabilities), continuous iteration through 2025 to 2026.

What the feature is: Intercom’s AI customer support agent. Fin answers customer questions autonomously by reading a company’s help content, past tickets, and knowledge base. Priced per resolution (charged only when Fin actually resolves a ticket without human handoff). Publicly shared metrics on resolution rates often exceed 50% for well configured deployments.

What wasn’t available before: Pre Fin, Intercom had chatbots (decision tree based) and human agent tools. Automated resolution without a human required extensive bot configuration and handled narrow use cases. Fin was the first AI agent inside Intercom that could resolve open ended customer questions at scale, without bot builders setting up every flow.

How the product improved:

- **Redefined the unit economics of customer support.** Per resolution pricing aligned cost with value, which most AI products still don’t.
- **Made autonomous support credible at scale.** Intercom’s published customer outcomes (resolution rates, cost savings) gave the category real benchmarks to point to.
- **Repositioned Intercom entirely.** From “customer messaging platform” to “AI first customer service,” including public company rebranding and strategy shifts.
- **Set the reference for AI agents in a professional domain.** Competitors (Zendesk AI, Ada, Decagon, Sierra, Forethought) all respond to Fin’s benchmarks.

Strategic significance for the project: Fin is arguably the single most successful traditional SaaS AI feature in the dataset. Every project needs a clear success story to anchor “what good looks like” patterns, and Fin is it.

26. Atlassian Rovo

Launched: April 2024.

What the feature is: Atlassian’s standalone AI work assistant. Connects to Atlassian products (Jira, Confluence) and third party tools (Google Drive, Slack, Microsoft 365), lets users search and ask questions across that content, and provides agents for specific tasks. Distinct from Atlassian Intelligence, which is AI embedded inside Jira and Confluence.

What wasn’t available before: Pre Rovo, Atlassian had Intelligence features embedded in each product (AI search in Confluence, smart suggestions in Jira). There was no separate, standalone AI product that worked across Atlassian plus third party content, and no agent layer. Users who wanted cross tool AI used Glean or Dropbox Dash.

How the product improved:

- **Gave Atlassian an answer to Glean and Copilot.** The “AI across all your work tools” category now had an Atlassian product.
- **Extended Atlassian’s reach beyond Jira and Confluence.** Rovo connects to non Atlassian tools, positioning it as a standalone product category rather than an add on.
- **Introduced the agent layer.** Rovo Agents for specific tasks (research, writing, analysis) moved Atlassian into the agent wave.
- **Provided a bundled AI pitch for enterprise sales.** Existing Atlassian customers got a cross tool AI option without adding another vendor.

Strategic significance for the project: Rovo sets up a cross dataset comparison with Glean and Dropbox Dash. All three are chasing the same “AI across your work” thesis from different starting points. Scoring them together reveals whether incumbent or native approaches win.

27. Airtable AI (AI Fields + Cobuilder)

Launched: March 2024 (AI Fields), 2024 (Cobuilder).

What the feature is: Airtable’s AI layer applied to structured data. AI Fields let users apply AI to any table column (summarize content, categorize entries, translate, extract information, generate text). Cobuilder lets non technical users build Airtable apps from natural language prompts. Together, they add AI across Airtable’s database and app building surfaces.

What wasn’t available before: Pre AI Fields, Airtable had scripting and automations, but applying AI to column data required custom scripts or third party integrations. There was no one click “add an AI column” flow. Pre Cobuilder, building an Airtable app meant manually configuring tables, views, and interfaces.

How the product improved:

- **Introduced AI for structured data.** Categorizing 10,000 support tickets or translating a product catalog became trivial.
- **Lowered the Airtable learning curve.** Cobuilder meant less technical users could start with a working app and edit from there.
- **Extended Airtable’s positioning against competitors.** Notion, Coda, and Smartsheet all had AI bets; Airtable’s structured data approach was its differentiation.
- **Opened new enterprise use cases.** AI on structured data unlocked operations, HR, and marketing workflows that were previously too manual.

Strategic significance for the project: Airtable AI is the dataset’s clearest “AI for structured data” entry. Most AI product discussion focuses on chat interfaces; structured data AI has underexplored adoption dynamics, making this one of the more original entries.

28. Miro AI (Intelligent Canvas)

Launched: 2024 (Intelligent Canvas rollout).

What the feature is: Miro's umbrella AI layer for visual collaboration. Features include auto clustering sticky notes, summarization of whiteboard content, AI generated diagrams from prompts, AI generated mind maps, presentation generation from Miro boards, and AI assisted brainstorming. Shipped as many small features rather than one big bet.

What wasn't available before: Pre Miro AI, Miro was a manual whiteboard tool. Users organized sticky notes by hand, created diagrams from scratch, and synthesized workshop outputs manually. AI features existed in beta form but weren't integrated as a product layer.

How the product improved:

- **Automated workshop cleanup.** After a remote brainstorming session, Miro AI could cluster and summarize outputs that previously took a facilitator hours.
- **Lowered the diagram creation floor.** Users could describe a diagram and get a reasonable starting point, instead of starting blank.
- **Positioned Miro against competitors.** FigJam, Mural, Lucidspark, and even Figma itself were shipping AI in collaboration tools.
- **Added new creative surfaces.** Mind map generation, presentation export, and AI assist across the canvas expanded Miro's use cases.

Strategic significance for the project: Miro is the dataset's weakest traditional SaaS entry due to scattered AI and thin public evidence. Its survival scoring will need to be more qualitative than most, and this is worth acknowledging in the judgment calls appendix.

29. Grammarly Generative AI Suite

Launched: April 2023 (initial GenAI rollout), with continuous expansion through 2024 to 2026.

What the feature is: Grammarly’s generative AI layer, including rewrite, tone adjustment, compose from prompt, and ideate. Works across 500,000 plus apps and sites where Grammarly operates (Gmail, Google Docs, LinkedIn, Slack, etc.). Marketed as Grammarly transitioning from “grammar checker” to “AI writing assistant.”

What wasn’t available before: Pre GenAI Suite, Grammarly was a correction tool. It fixed errors, improved clarity, and suggested word choices. It did not generate new text. Users who wanted AI composition went to ChatGPT and pasted outputs back into their documents. Grammarly’s risk was becoming a commodity feature as writing AI became baseline.

How the product improved:

- **Defended Grammarly’s market against ChatGPT encroachment.** Without GenAI features, Grammarly’s value would have collapsed as users increasingly composed directly in ChatGPT and bypassed correction.
- **Expanded use cases beyond correction.** Generate, rewrite, and tone shift features gave users a reason to stay in Grammarly during composition, not just after.
- **Leveraged Grammarly’s unique distribution advantage.** Because Grammarly operates inside the apps users already write in, GenAI was embedded in workflow without switching tools, which ChatGPT couldn’t match.
- **Enabled upsell to GrammarlyGO and higher tier plans.** GenAI features gated premium tiers, supporting the business model shift.

Strategic significance for the project: Grammarly is the dataset’s only “successful re platforming” story. Scoring whether its users actually adopted GenAI (or kept using Grammarly just for grammar while composing in ChatGPT) is the strategic question at the heart of its survival.

30. Dropbox Dash

Launched: Beta in 2023, general availability December 2024 (Dash for Business), Spring 2025 release, Fall 2025 release.

What the feature is: Dropbox’s universal AI search and knowledge management product. Connects to Google Drive, Slack, Microsoft 365, Notion, Canva, Jira, and other apps. Lets users search and ask questions across all connected apps with AI generated answers, citations, and action links. Includes Stacks (shared team workspaces) and an MCP server that exposes Dash content to Claude, Cursor, and other AI clients. Available as a standalone Dash app or embedded in the familiar Dropbox interface.

What wasn’t available before: Pre Dash, Dropbox was a file sync and share company. Customers stored files in Dropbox but searched, asked questions, and worked across other apps (Google Drive, Slack, Notion) separately. There was no AI layer across content. Customers who wanted cross app AI search used Glean or Microsoft Copilot.

How the product improved:

- **Transformed Dropbox’s product positioning.** From “the folder for your files” to “the AI layer for all your work content.” This reposition is the entire 2025 to 2026 strategic bet.
- **Reclaimed existing Dropbox customers’ attention.** Dropbox’s core product had been commoditized by Google Drive and iCloud. Dash gave customers a reason to re engage.
- **Positioned Dropbox against Glean, Microsoft Copilot, ChatGPT with connectors.** The “AI that knows your work content” pitch now had a real product.
- **Drove measurable business impact.** Dropbox publicly raised 2025 revenue guidance citing Dash adoption, which is unusually direct financial evidence for any AI feature in the dataset.

Strategic significance for the project: Dash is the clearest “traditional SaaS bet the company on AI” story in the dataset. Scoring whether Dash adoption actually justifies Dropbox’s corporate pivot is the single highest stakes question for a company at this size. Direct comparison with Glean sets up one of the project’s most informative cross entries.

How to use this doc during scoring

For each feature, when you sit down to score it, re read the relevant entry here first. Pay attention to:

Claim checking. Does the “how the product improved” story match what you’re finding in current evidence (G2 reviews, Reddit, changelogs)? If the feature was supposed to “compress drafting time for lawyers” and the reviews say “it hallucinates citations,” note the gap. That gap is where your insight angle lives.

Scope confirmation. Does the scope described here still match what you’re scoring? If Granola Spaces has since absorbed other features, or if Slack AI has added new sub features, update the scope.

Insight angle prompts. Each entry implies a specific question the feature hasn’t answered publicly. For ChatGPT Memory, it’s “did it actually change behavior or just sit as a toggle.” For Figma AI, it’s “what happened after the pause.” Those questions are the angles that make your writeup worth reading.

Version history

v1.0, April 2026. Initial creation during Day 1 of project scoping. 30 features across 30 companies. Expect corrections as scoring evidence surfaces gaps.