

After the Launch Party, Scoring Rubric

The methodology document for assigning an Adoption Score to each of the 30 AI features in the project.

Project thesis: Scoring 30 AI features from 30 companies (15 AI native, 15 traditional SaaS) to answer which features stuck roughly a year after launch, and what patterns separate the survivors from the shelved.

What this document is: The rubric. The instrument every feature in the dataset will be measured against. Built so that a second researcher could pick it up, score the same feature independently, and land within roughly a point of where the project author would land.

Version: v1.0, April 2026.

Section 1. Overview

Every feature in the dataset receives an **Adoption Score from 0 to 10**, computed as a weighted average of five signals.

Signal	Weight	What it captures
Usage evidence	25%	Direct or proxy evidence that people actually use the feature
User sentiment	20%	What users say when they do talk about the feature
Company conviction	20%	Whether the company still acts like the feature matters
Competitive response	15%	Whether peers shipped something similar within roughly a year

Signal	Weight	What it captures
Strategic signal	20%	Whether pricing, positioning, and leadership voice still back the feature

Why these five. The split is deliberate. Forty five percent of the score lives on the user side (usage plus sentiment), thirty five percent on the company side (conviction plus strategic), and fifteen percent on the market (competitive response). A project asking what survived needs to weight the people using the feature more heavily than the people shipping it. The original draft of this rubric leaned the other way, and we corrected for it.

What the score means in plain language:

- **0 to 3.** Shelved, deprecated, ignored, or quietly dying. Few users, weak sentiment, fading from the company narrative.
- **4 to 6.** Alive but unremarkable. Some users, mixed sentiment, still in the product but not a centerpiece.
- **7 to 8.** Genuinely sticky. Real usage, real positive sentiment, still strategically backed, often imitated by competitors.
- **9 to 10.** Category defining. The feature changed user expectations across the space. Rare.

Section 2. The Five Signals, Defined

Each signal below has four parts. A definition. A list of where to look. Anchor descriptions for what a 2, 5, and 8 look like. And a rule for what to do when the signal cannot be found.

Signal 1. Usage Evidence (25%)

Definition. Direct or proxy evidence that real users are actually using the feature, not just that the company is shipping it. This is the single most important question the project is trying to answer, which is why it carries the heaviest weight.

Where to look.

- Reddit threads and subreddit activity related to the product. Look for recent posts (within the last three months of the scoring date) that mention the feature.
- YouTube tutorials and walkthroughs. Search for the feature name. Check view counts and recency.
- Third party tools, integrations, browser extensions, or plugins built specifically around the feature. Their existence implies real usage.
- Public usage statistics published by the company. Treat these with mild skepticism (companies pick favorable numbers) but include them.
- Support volume signals. Help center articles updated recently, community forum questions, Stack Overflow tags.
- Power user content. Newsletters, blog posts, and Twitter/X threads from heavy users describing how they use the feature.

Anchors.

- **Score 2.** Almost no organic discussion. Zero or near zero Reddit threads in the last quarter. No tutorial videos worth mentioning, or only the company's own promotional video. No third party tools. The feature exists, but the silence is loud.
- **Score 5.** Modest activity. A handful of Reddit threads, some tutorial videos with low to mid view counts, occasional power user posts. Enough to confirm the feature has real users, not enough to suggest it is widely loved or widely adopted.
- **Score 8.** Robust ecosystem. Active and recent Reddit discussion, multiple tutorial videos with strong view counts, third party tools or extensions, power user newsletters, and frequent how to questions. People are visibly building habits around the feature.

Missing data rule. Score a 2. Silence is data here. If a feature has launched and a year later there is essentially no organic conversation about it, that absence is itself the answer.

Signal 2. User Sentiment (20%)

Definition. When users do talk about the feature, what is the tone and substance of what they say. This is distinct from usage evidence, which measures whether they talk at all. Sentiment measures what they say once they do.

Where to look.

- Reddit threads and comments. Read the actual content, not just the count.
- G2, Capterra, and Trustpilot reviews that specifically mention the feature.
- Twitter/X mentions and quote tweets about the feature.
- Product forum threads on the company's own community site.
- App store reviews where applicable.
- YouTube comments on the company's launch video and on tutorial videos.

Anchors.

- **Score 2.** Predominantly negative. Common themes include “it does not work,” “it is wrong half the time,” “I turned it off,” “it slowed down the product.” Or a pattern of users actively requesting the feature be removed or made optional.
- **Score 5.** Mixed. Some users find it useful, some are frustrated, and a meaningful portion are indifferent or unaware. Common phrasing: “it is fine,” “it is okay for X but not Y,” “I use it sometimes.”
- **Score 8.** Predominantly positive. Users describe specific use cases where the feature saves them time or solves a real problem. Common themes include “I use this every day,” “I cannot go back,” or specific gratitude. Few complaints, and the complaints are usually about wanting more, not wanting less.

Missing data rule. Score a 2. Same reasoning as usage evidence. If nobody is saying anything about the feature a year after launch, that is itself a strong negative signal.

Signal 3. Company Conviction (20%)

Definition. Whether the company still behaves as if the feature matters. This is a merger of the original “prominence decay” and “changelog persistence” signals from the first draft, because both were measuring the same underlying thing: does the company still believe in this.

Where to look.

- The product itself. Is the feature still surfaced in onboarding flows, dashboards, default views, and primary navigation, or has it been buried in a submenu or removed.
- The marketing site. Does the homepage, the features page, or a dedicated landing page still showcase it.
- Recent changelogs and release notes. Has the feature received updates in the last six months, or has it gone silent.
- The pricing or plans page. Is the feature listed, and is it positioned prominently or buried.
- The company blog. Recent posts about the feature, customer stories featuring it, or technical deep dives.

Anchors.

- **Score 2.** Buried or gone. Removed from the homepage, no recent changelog mentions, hidden in submenus, or quietly deprecated. The product team has effectively moved on even if the feature still technically exists.
- **Score 5.** Still present, not centered. The feature is in the product and on some pages but is no longer a marketing centerpiece. Updates are infrequent. It is being maintained, not invested in.
- **Score 8.** Visibly invested in. The feature appears prominently on the marketing site, gets regular changelog updates, is featured in customer stories, and shows up in onboarding. The company is actively building on it.

Missing data rule. Skip and reweight. This is rare for company conviction (public companies almost always have findable signals here), but if it happens, drop the signal and rescale the remaining four to sum to 100%.

Signal 4. Competitive Response (15%)

Definition. Whether other companies in the same space shipped something meaningfully similar within roughly a year of the original launch. A feature copied by competitors is a feature the market has validated. A feature nobody bothered to copy is a feature the market quietly judged and dismissed.

Where to look.

- Direct competitors' product pages, changelogs, and launch announcements within twelve months following the original feature launch.
- Industry roundup articles, "X versus Y" comparison posts, and analyst reports that group similar features.
- Reddit and Twitter/X discussion that mentions multiple products having the same capability.
- Open source clones, if applicable. The existence of an open source version of a proprietary feature is a strong validation signal.

Anchors.

- **Score 2.** No imitation. Twelve months later, no major competitor has shipped a comparable feature. The feature stands alone, which usually means the market does not believe in it.
- **Score 5.** Partial imitation. One or two competitors shipped something similar, often as a smaller bet or a marketing checkbox rather than a serious investment.
- **Score 8.** Widespread imitation. Most or all major competitors shipped a comparable feature within twelve months. The category has shifted to treat this capability as table stakes.

Missing data rule. Skip and reweight. If the competitive landscape genuinely cannot be assessed (rare, but possible for niche verticals), drop the signal and rescale.

Signal 5. Strategic Signal (20%)

Definition. Whether the company’s pricing, positioning, and leadership voice still back the feature. This is a merger of the original “pricing signal” and “PM/exec voice” categories, because both measure the same thing: is the leadership of the company still betting on this. Pricing alone has a coverage problem (bundled features have no signal), and exec voice alone is too narrow. Together they form a stronger composite.

Where to look.

- Pricing page. Is the feature gated behind a paid tier, included in a premium plan, metered, or bundled freely. Has its pricing position changed in the last twelve months.
- Earnings calls and investor materials, for public companies. Searches like “company name earnings call feature name.”
- Podcast appearances by the CEO, CPO, or product leads. Look for mentions of the feature within the last six to twelve months.
- LinkedIn posts from leadership.
- Conference keynotes and product launch events.
- Press interviews and feature stories.

Anchors.

- **Score 2.** Strategic withdrawal. The feature has been removed from premium tiers, made free as an afterthought, or dropped from leadership talking points. Or both. The company is signaling, through its actions, that this is no longer a strategic bet.
- **Score 5.** Still bet on, not centered. The feature is priced reasonably (paid tier or premium add on) and gets occasional mentions from leadership, but is no longer a headline talking point. It is part of the portfolio, not the centerpiece.
- **Score 8.** Active strategic priority. The feature anchors a paid tier or has clear pricing power, and leadership talks about it frequently in earnings, podcasts, and keynotes. The company is publicly betting on it as a growth driver.

Missing data rule. Skip and reweight. Strategic signals are often unavailable for private companies that do not disclose pricing publicly or whose leadership has a low public profile. Absence here usually reflects communication style rather than feature fate, so penalizing the score is unfair.

Section 3. Time Horizon

Default scoring window. Twelve months post launch. For a feature launched in October 2024, signals are evaluated as of October 2025.

Exception. For features launched after April 2025, a full twelve month window does not exist as of the project date. These features are scored “as of today” with the actual elapsed time noted in the scoring sheet (for example, “scored at eight months post launch, April 2026”).

How to note it. Each feature in the scoring sheet has a column for “scoring window” with one of two values: “12 months” or “X months, partial window.” The methodology section of the final report flags this transparently.

Why this matters. A reader can see exactly which features got the full treatment and which got a shorter window, and can decide for themselves how that affects the conclusions. Acknowledging limitations openly is part of what separates research from punditry.

Section 4. Final Score Calculation

The formula.

$$\text{Adoption Score} = (\text{Usage} \times 0.25) + (\text{Sentiment} \times 0.20) + (\text{Conviction} \times 0.20) + (\text{Competitive} \times 0.15) + (\text{Strategic} \times 0.20)$$

When a signal is skipped (per the missing data rules above). Drop the missing signal and rescale the remaining weights to sum to 1.0. Example: if strategic signal is unavailable, the remaining four weights of 0.25, 0.20, 0.20, and 0.15 sum to 0.80. Divide each by 0.80 to get the rescaled weights of 0.3125, 0.25, 0.25, and 0.1875. Apply those.

Worked example. Notion AI Q&A, scored at twelve months post launch (November 2024).

Signal	Raw score	Weight	Contribution
Usage evidence	6.5	0.25	1.625
User sentiment	5	0.20	1.00
Company conviction	7	0.20	1.40
Competitive response	8	0.15	1.20
Strategic signal	6	0.20	1.20
Total			6.43

Notion AI Q&A Adoption Score: 6.4. Alive, modestly sticky, well imitated, but with mixed sentiment dragging it down. The score lines up with the qualitative read: it is a feature that works but does not yet feel essential.

Rounding convention. Final Adoption Scores are reported to one decimal place. Component signal scores are integers from 0 to 10. Half points (for example, 6.5 in the worked example above) are allowed for component scores when the evidence genuinely sits between two anchors, but should be used sparingly.

Section 5. Scoring Discipline

A rubric is only as good as the consistency with which it is applied. These are the practices that protect the project from scoring drift across 30 features.

Calibration check at the end. After all 30 features are scored, return to the first 5 and rescore them blind (without looking at the original scores). If any rescore differs from the original by more than one full point, investigate. The standard you applied in week one should match the standard you applied in week three. Drift is the silent enemy of comparability.

Document every judgment call. The scoring sheet has a “notes” column for each feature. Use it. Any time a score is between anchors, any time a signal is skipped, any time a tough call gets made, write a one line note explaining why. This is what makes the scoring auditable later.

Score one signal across all features before moving to the next, if possible. Scoring all 30 features on usage evidence first, then all 30 on sentiment, and so on, produces more consistent scores than scoring all five signals for one feature, then moving to the next. Working signal by signal forces direct comparison and surfaces calibration issues faster.

Resist the urge to round to a story. A feature you expected to be a hit might score a 5. A feature you expected to fail might score a 7. When that happens, trust the rubric, not the prediction. The whole point of the methodology is to let the evidence override the priors. The surprises are where the interesting findings live.

Cite as you go. For every component score, drop a link or a screenshot into the notes. Future you, reading the project six months from now, will not remember why a particular score landed where it did. Sources locked in at scoring time make the project defensible and reusable.

End of rubric. Day 2 deliverable, complete.